

CS 4530 Software Engineering

Module 13: Principles and Patterns of Cloud Infrastructure

Learning objectives for this lesson

By the end of this lesson, you should be able to...

- Describe what “cloud” computing is
- Understand the role of virtual machines and containers in cloud computing
- Deploy a web app to the cloud


How to deploy web apps?

What we need:

- ⦿ A server that can run our application
- ⦿ A network that is configured to route requests from an address to that server


Questions to think about:

- ⦿ What software do we need to run besides our application code?
- ⦿ Where does this server come from?
- ⦿ Who else gets to use this server?
- ⦿ Who maintains the server and software?



Many apps rely on common infrastructure

- Content delivery network: caches static content “at the edge” (e.g. cloudflare, Akamai)
- Web servers: Speak HTTP, serve static content, load balance between app servers (e.g. haproxy, traefik)
- App servers: Runs our application
- Misc services: Logging, monitoring, firewall
- Database servers: Persistent data



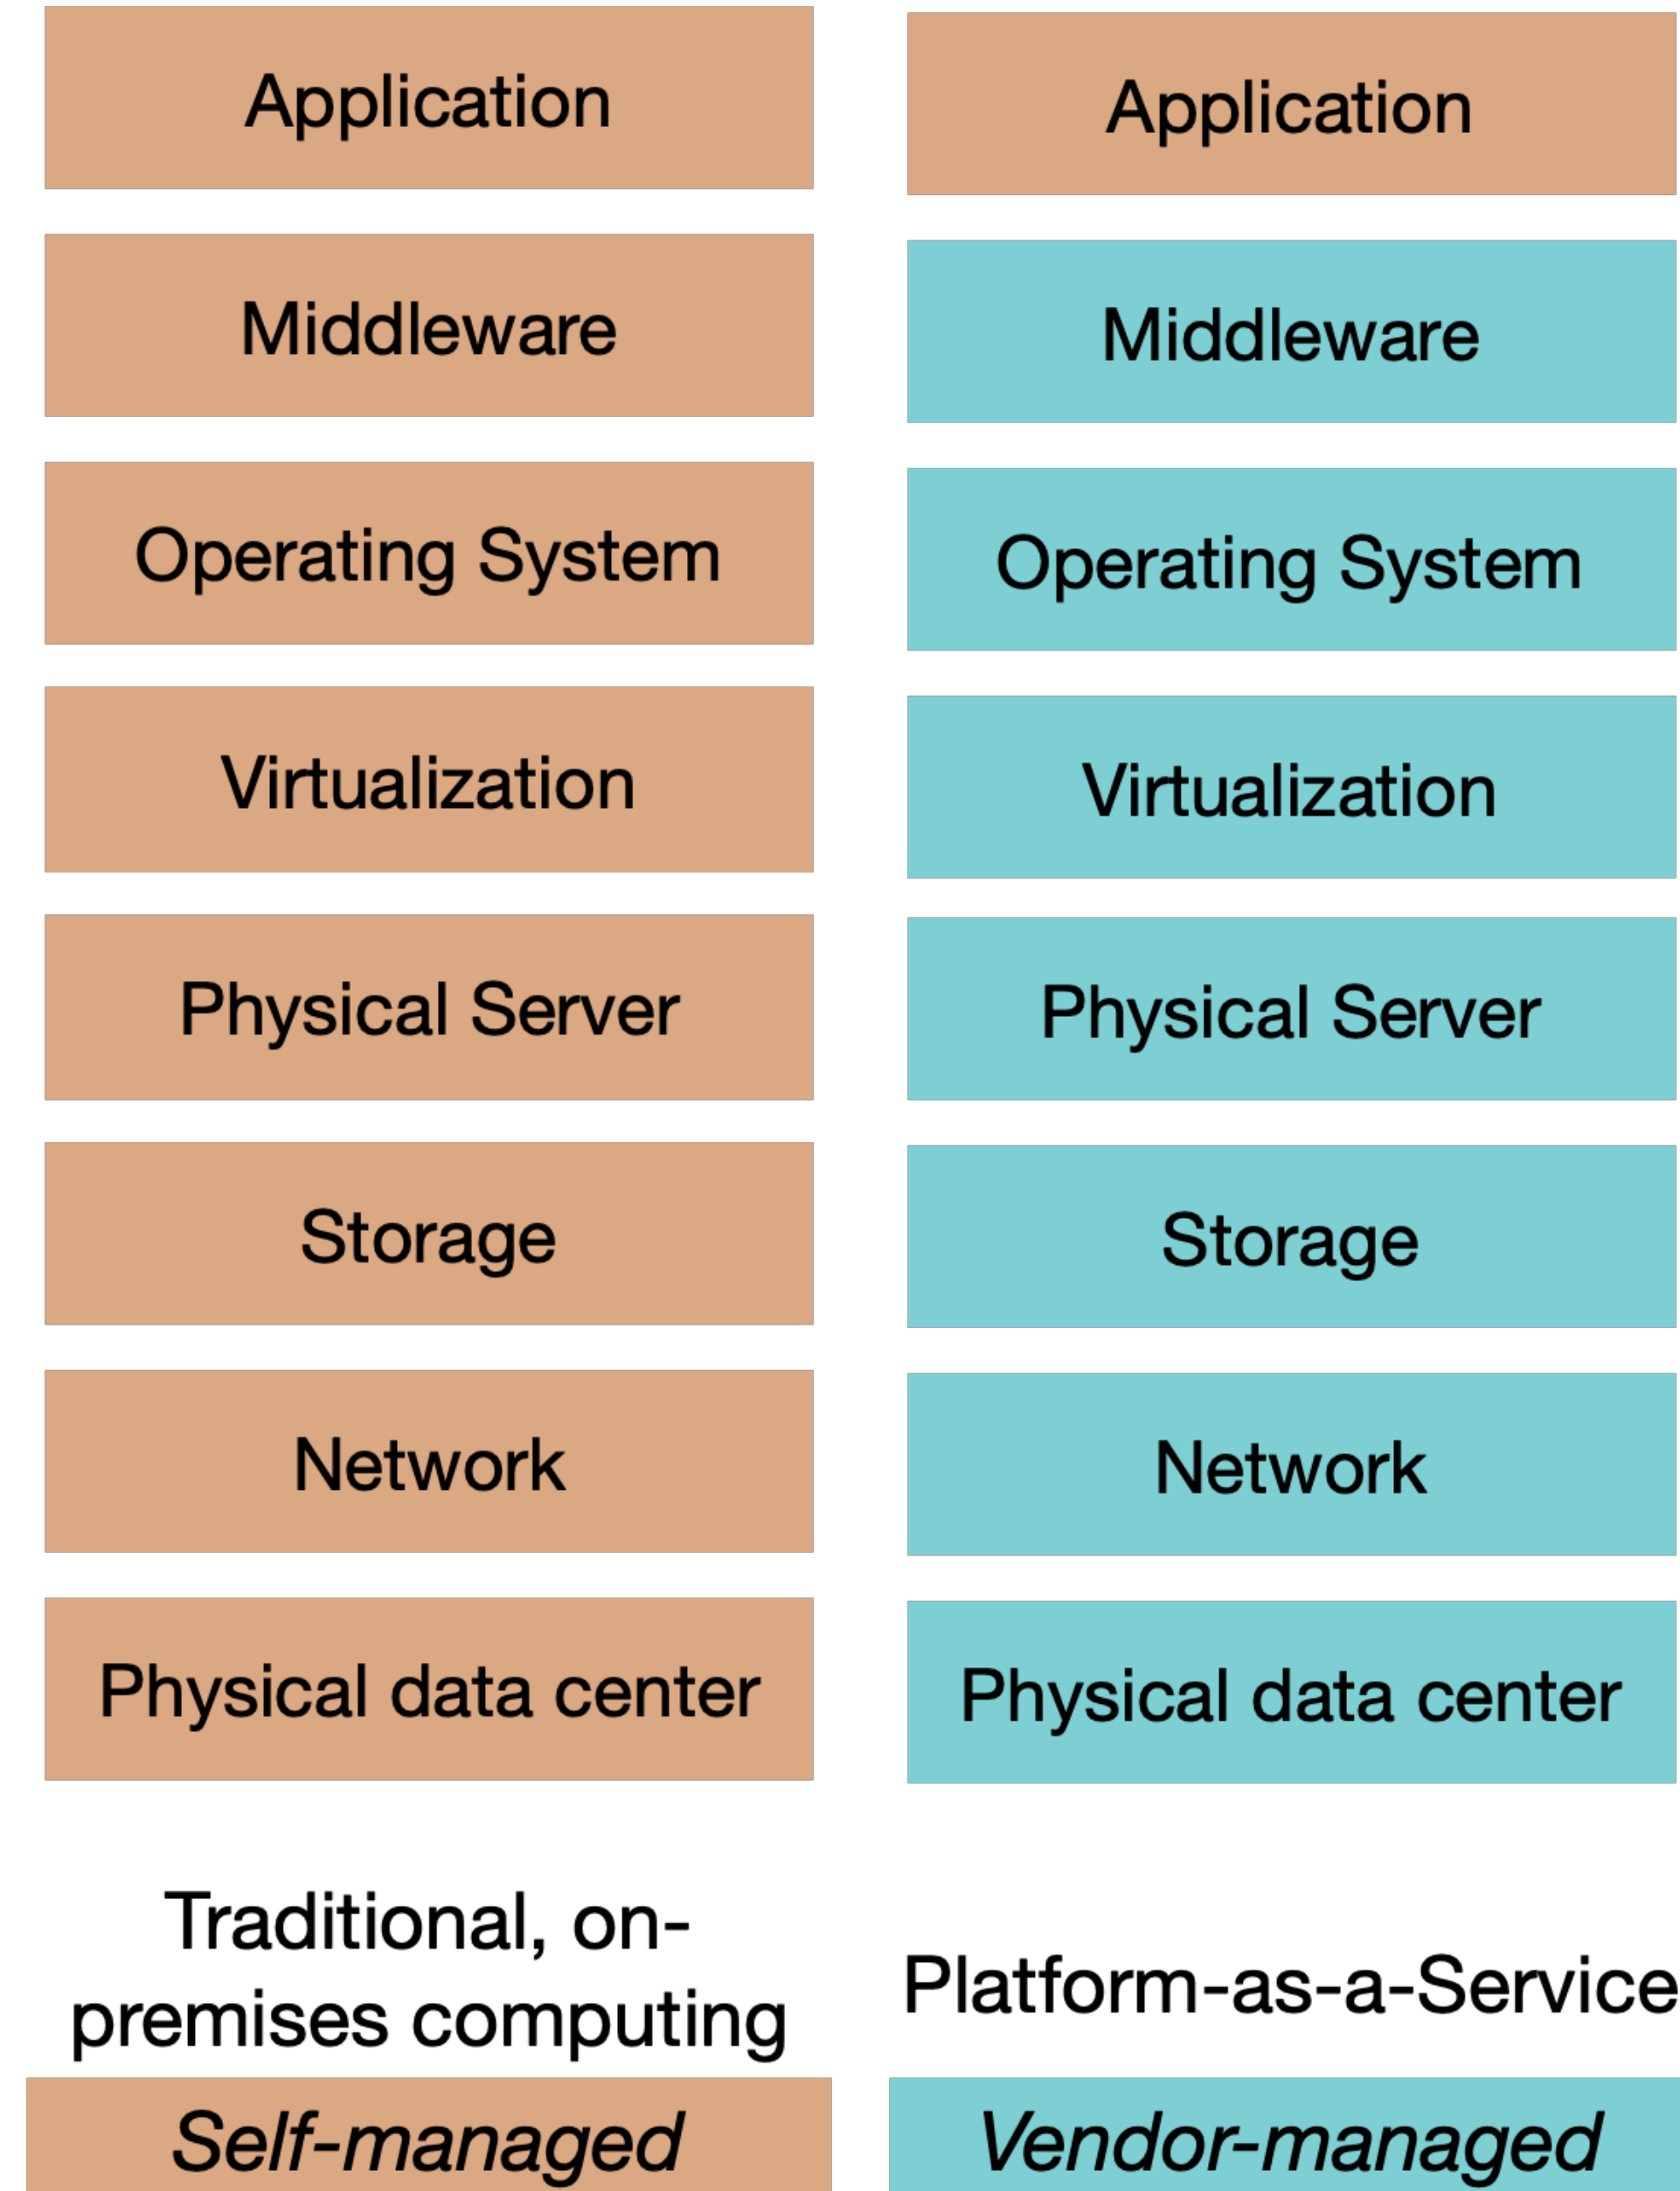
What is a cloud infrastructure?

Our apps run on a “tall stack” of dependencies

Traditionally this full stack is self-managed

Cloud providers offer products that manage parts of that stack for us:

- “Infrastructure as a service”
- “Platform as a service”
- “Software as a Service”



Cloud infrastructure creates economies of scale

At the physical level:

- ⦿ Multiple customers' physical machines in the same data center
- ⦿ Save on physical costs (centralize power, cooling, security, maintenance)

At the physical server level:

- ⦿ Multiple customers' virtual machines in the same physical machine
- ⦿ Save on resource costs (utilize marginal computing capacity)

At the application level:

- ⦿ Multiple customer's applications hosted in same virtual machine
- ⦿ Save on resource overhead (eliminate redundant infrastructure like OS)

Application

Middleware

Operating System

Virtualization

Physical Server

Storage

Network

Physical data center

*Multiple customers
could share each of
these tiers*

Cloud infrastructure scales elastically

“Traditional” computing infrastructure requires capital investment

- ⦿ “Scaling up” means buying more hardware, or maintaining excess capacity for when scale is needed
- ⦿ “Scaling down” means selling hardware, or powering it off

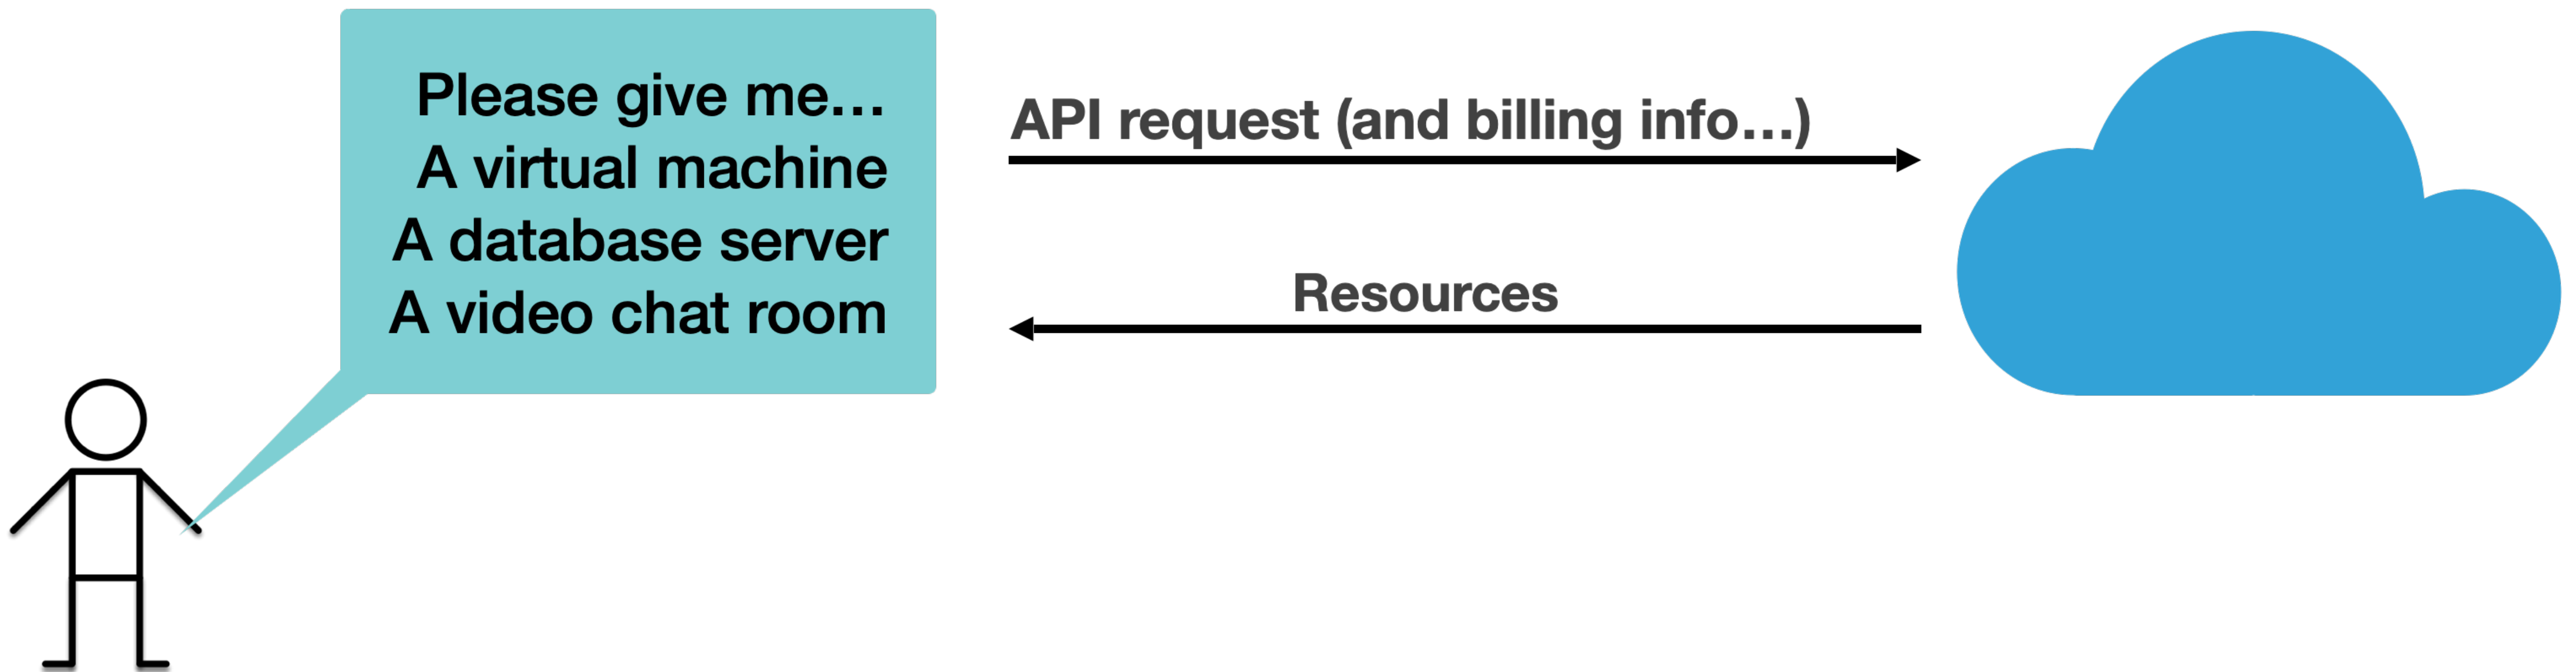
Cloud computing scales elastically:

- ⦿ “Scaling up” means allocating more shared resources
- ⦿ “Scaling down” means releasing resources into a pool
- ⦿ Billed on consumption (usually per-second, per-minute or per-hour)

Cloud Infra is on-demand access to resources

Vendor provides a service catalog of “X as a service” abstractions

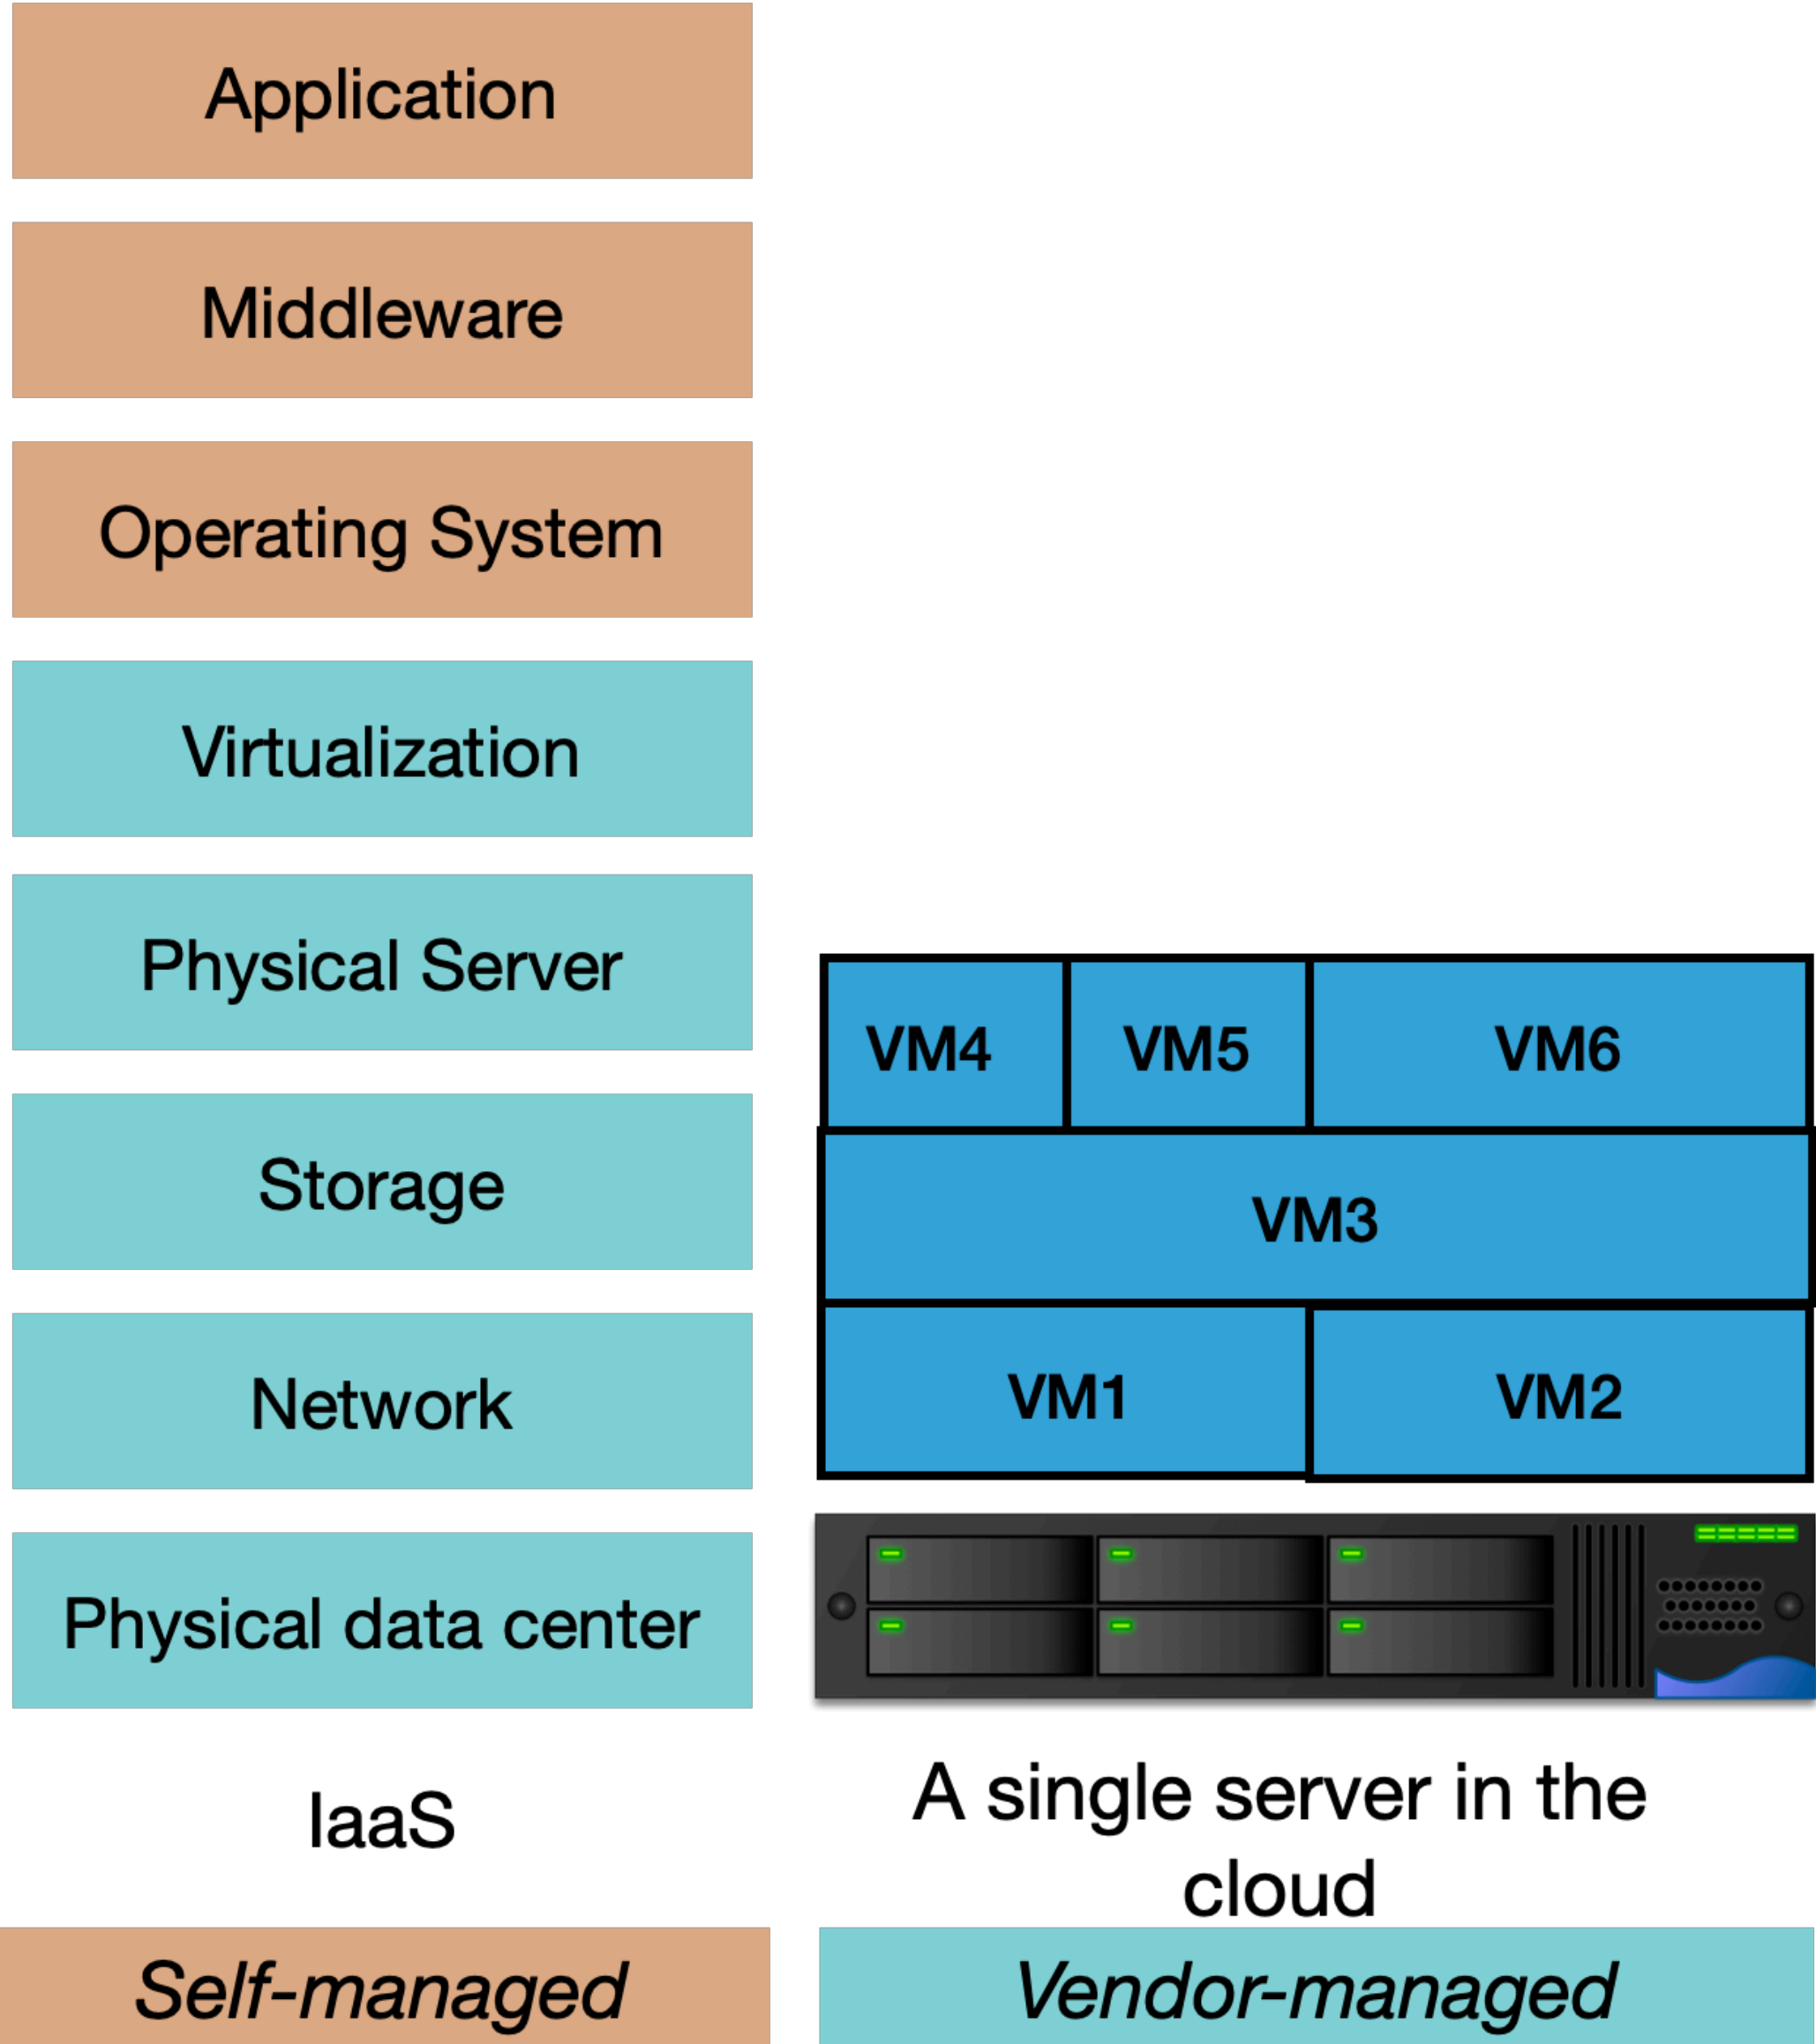
API allows us to provision resources on-demand



Infrastructure as a Service: Virtual Machines

Virtual machines:

- Virtualize a single large server into many smaller machines
- OS limits resource usage and guarantees quality per-VM
- Each VM in its own OS
- Examples: Amazon EC2, Google Compute Engine, Azure



Virtual Machines are a core abstraction

Multi-Tenancy

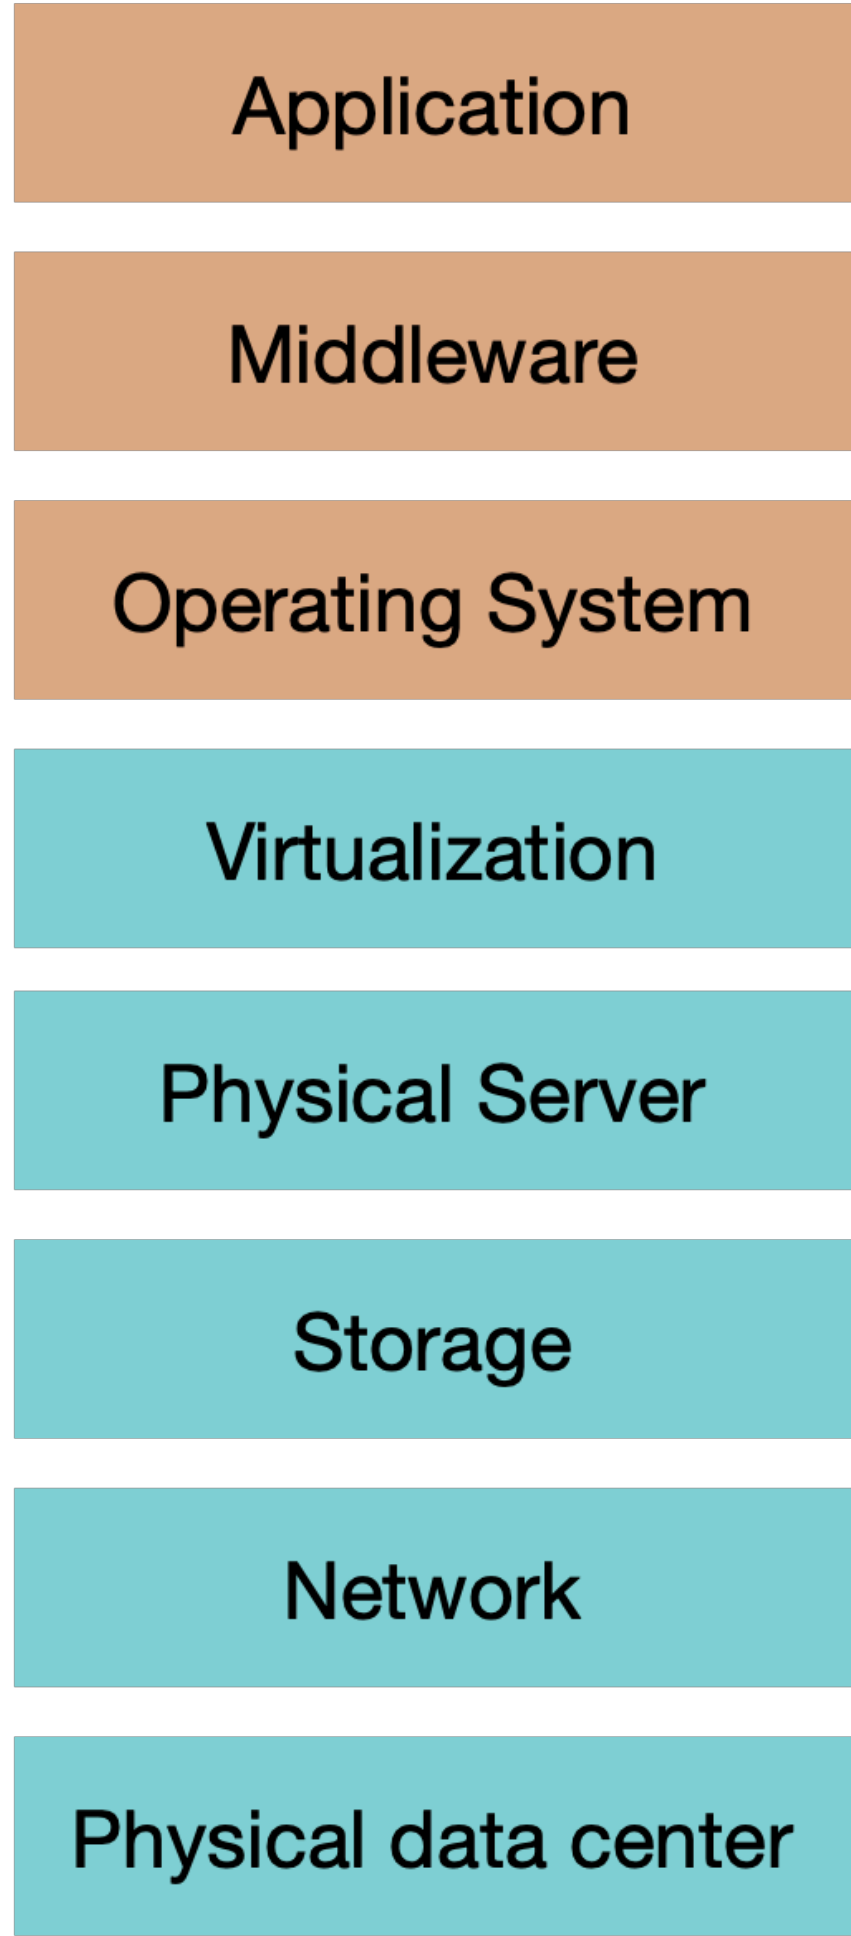
- Multiple customers sharing same physical machine, oblivious to each other

Decouples application from hardware

- virtualization service can provide “live migration”


Faster to provision and release

- VM v. physical machines == ~mins v. ~hours



IaaS

Self-managed



A single server in the cloud

Vendor-managed

Virtual Machines to Containers

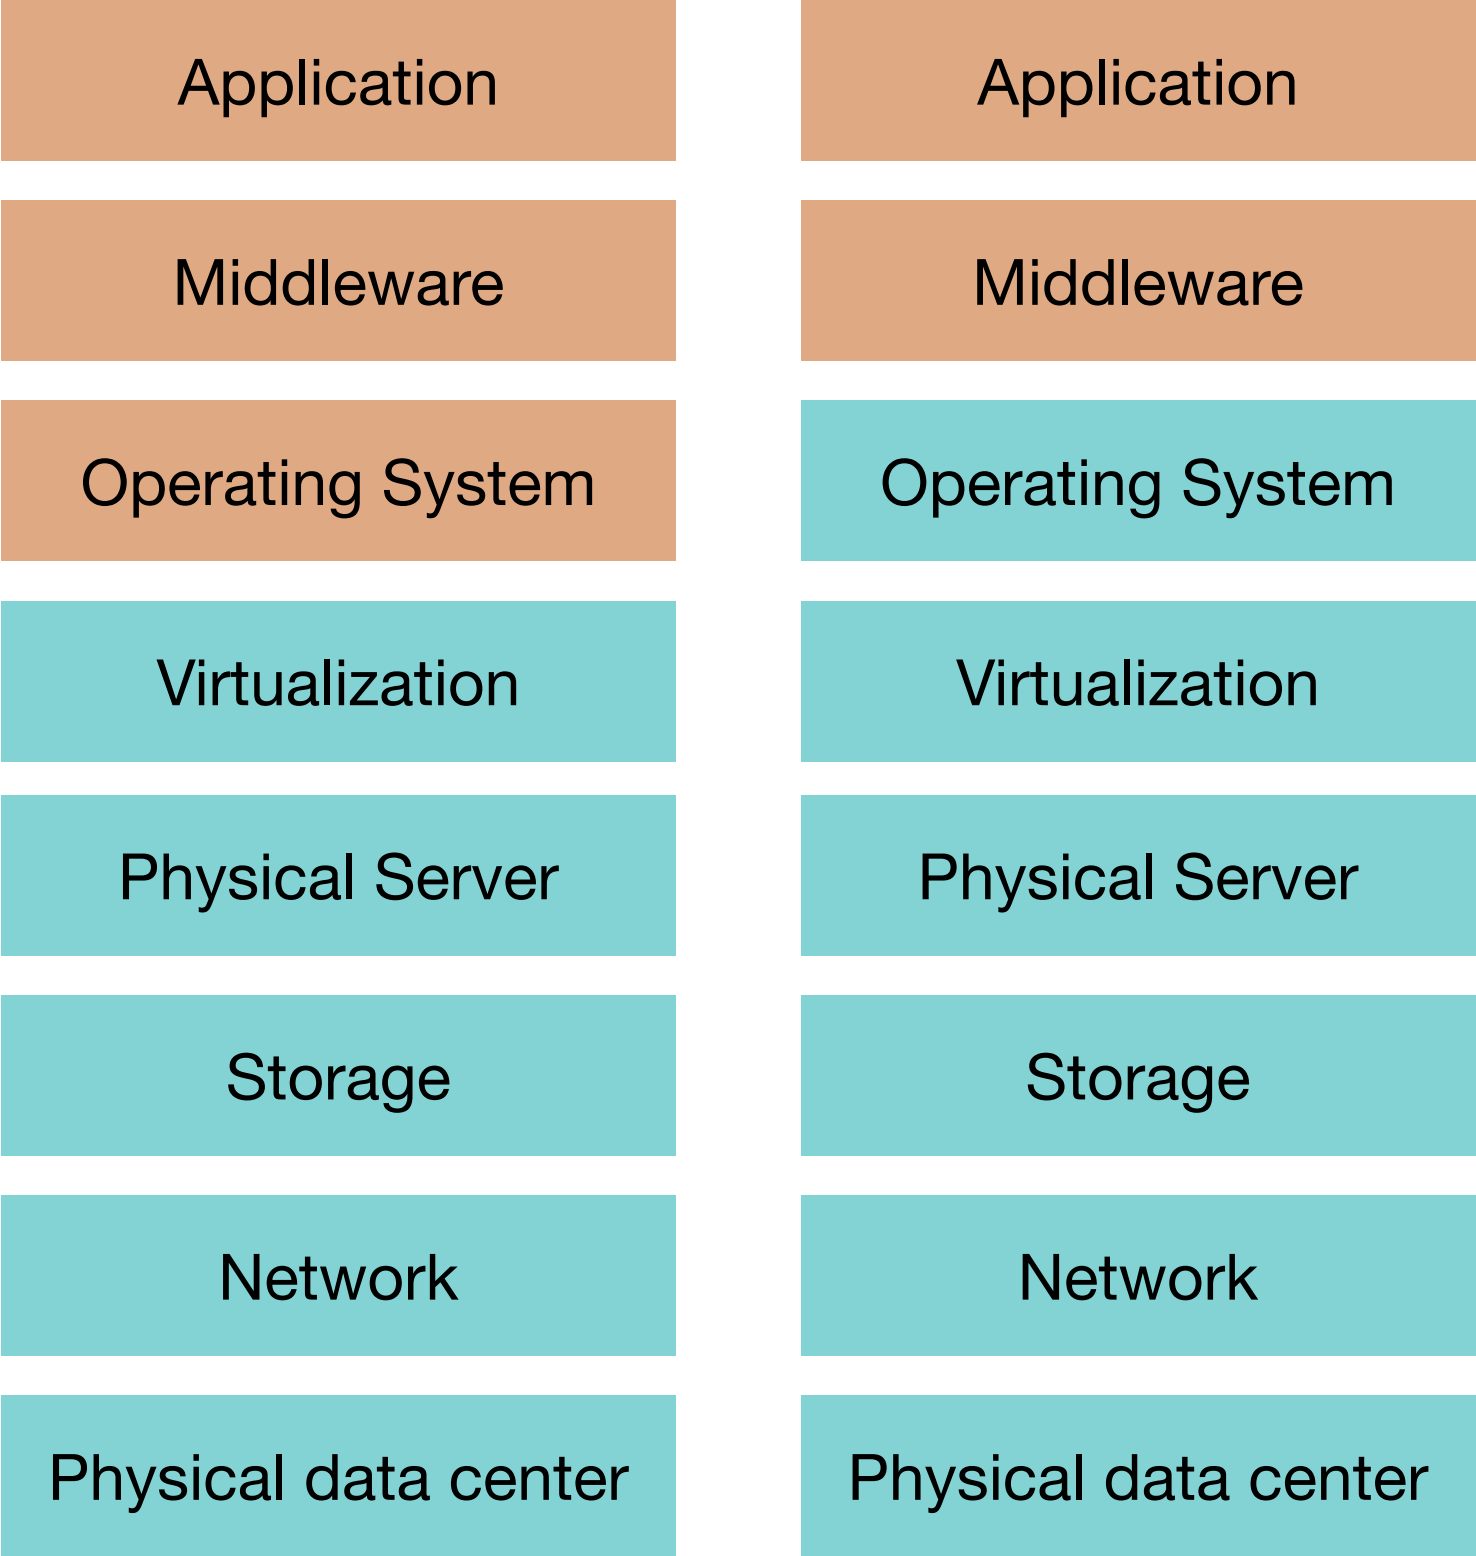
Each VM contains a full operating system

What if each application could run in the same (overall) operating system? Why have multiple copies?

Advantages to smaller apps:

Faster to copy (and hence provision)

Consume less storage at rest



VMs as a service

Containers as a
service

Self-managed

Vendor-managed

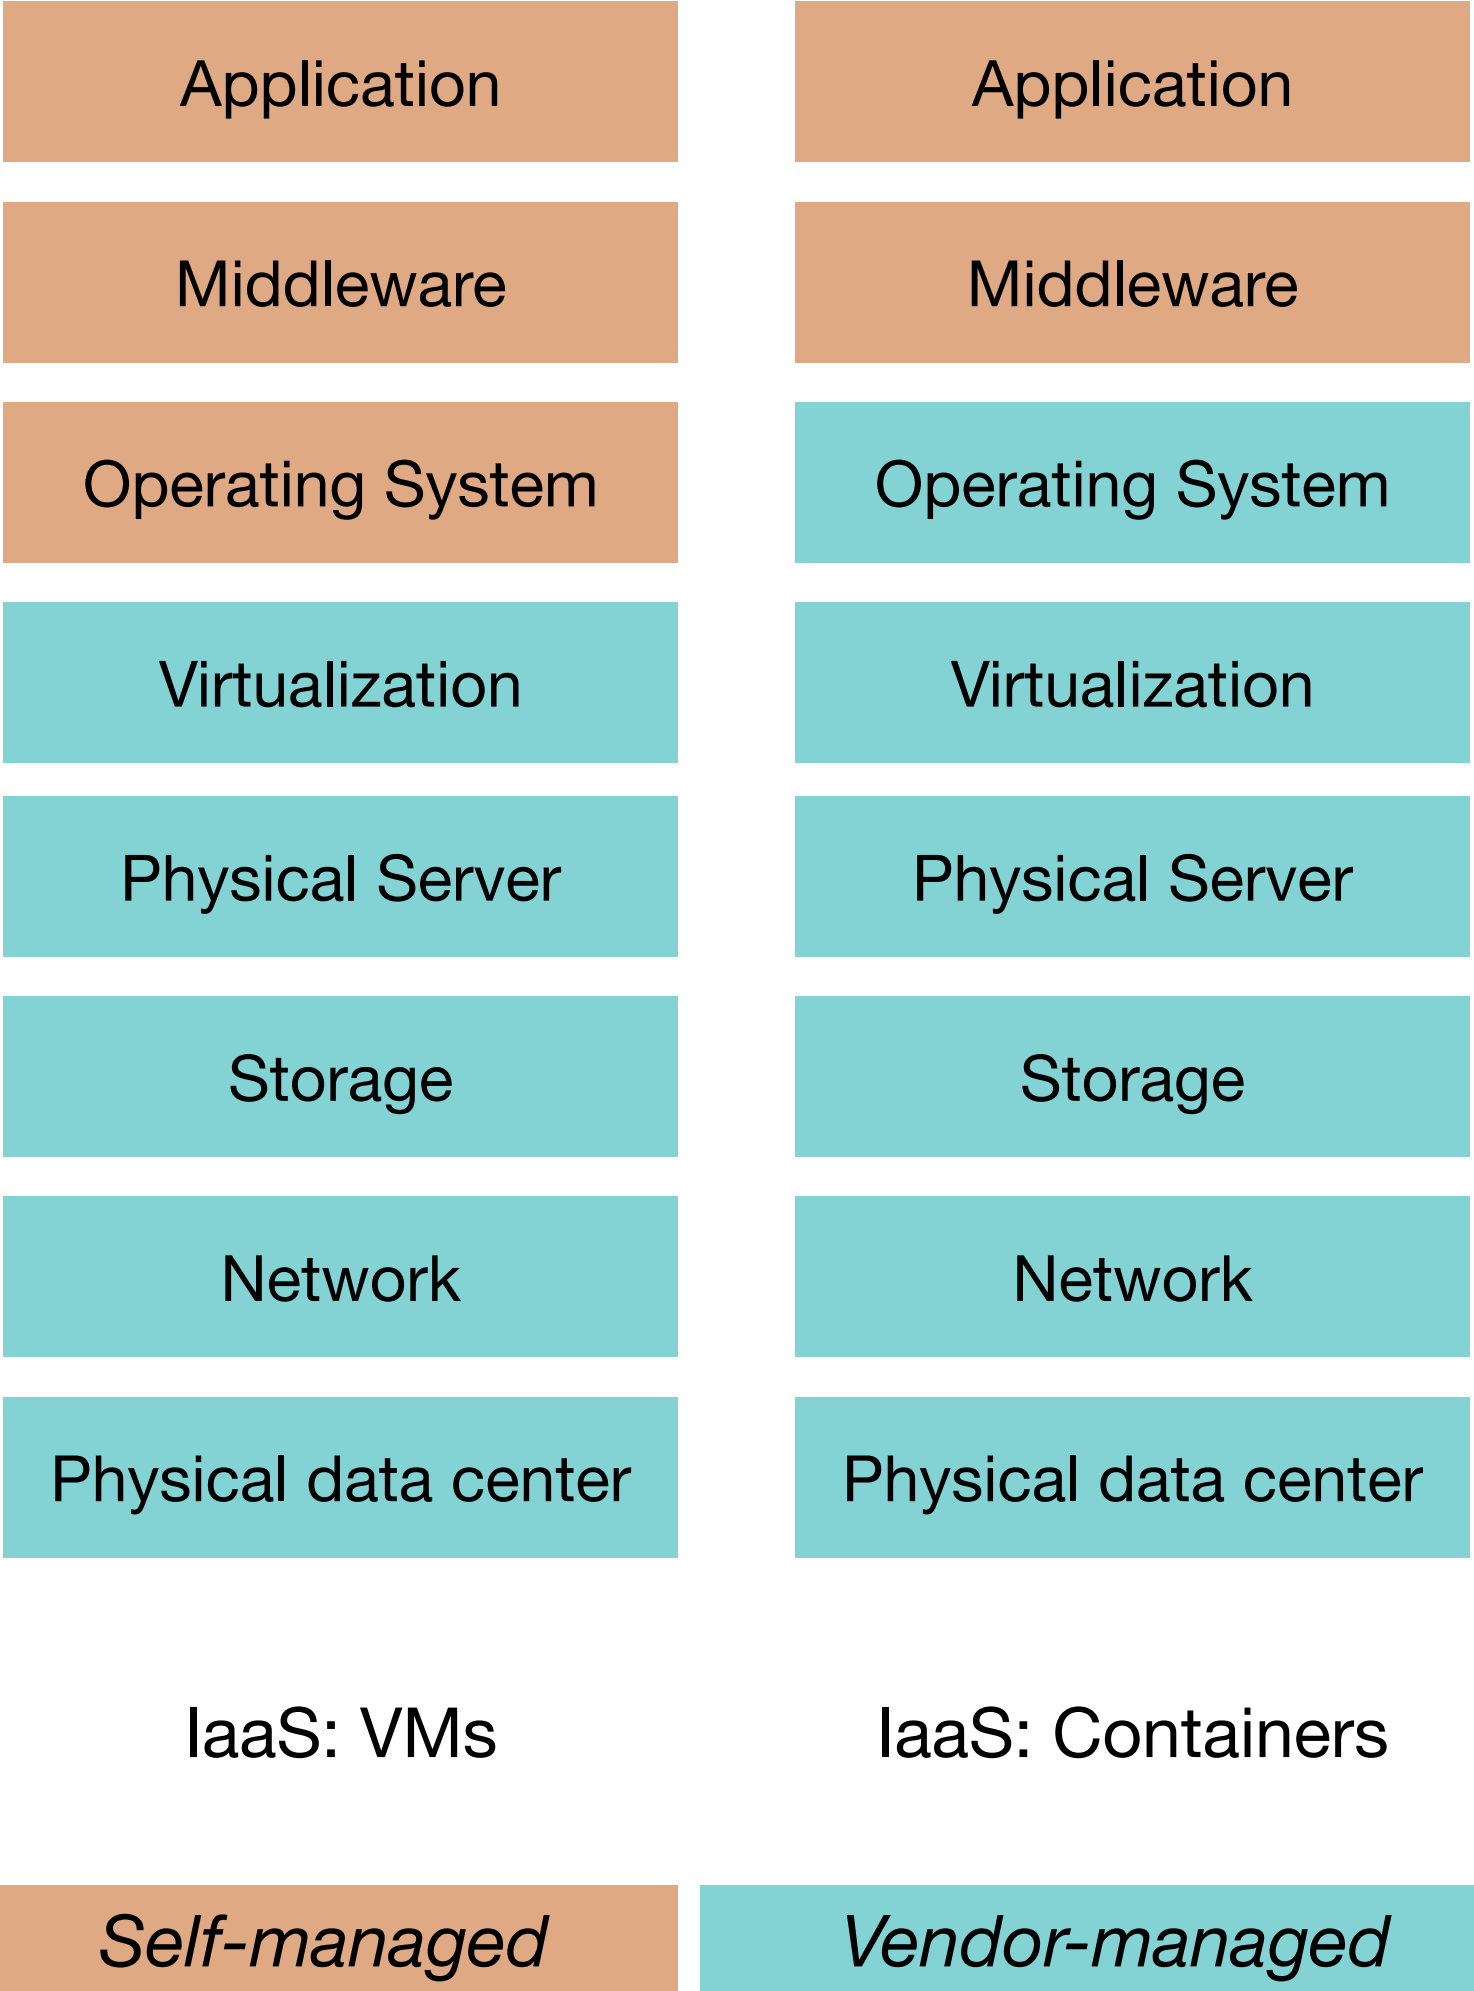
Infrastructure as a Service: Containers

Each application is encapsulated in a “lightweight container,” includes:

- System libraries (e.g. glibc)
- External dependencies (e.g. nodejs)

“Lightweight” in that container images are smaller than VM images - multi tenant containers run in the OS

Cloud providers offer “containers as a service”
(Amazon ECS Fargate, Azure Kubernetes, Google Kubernetes)

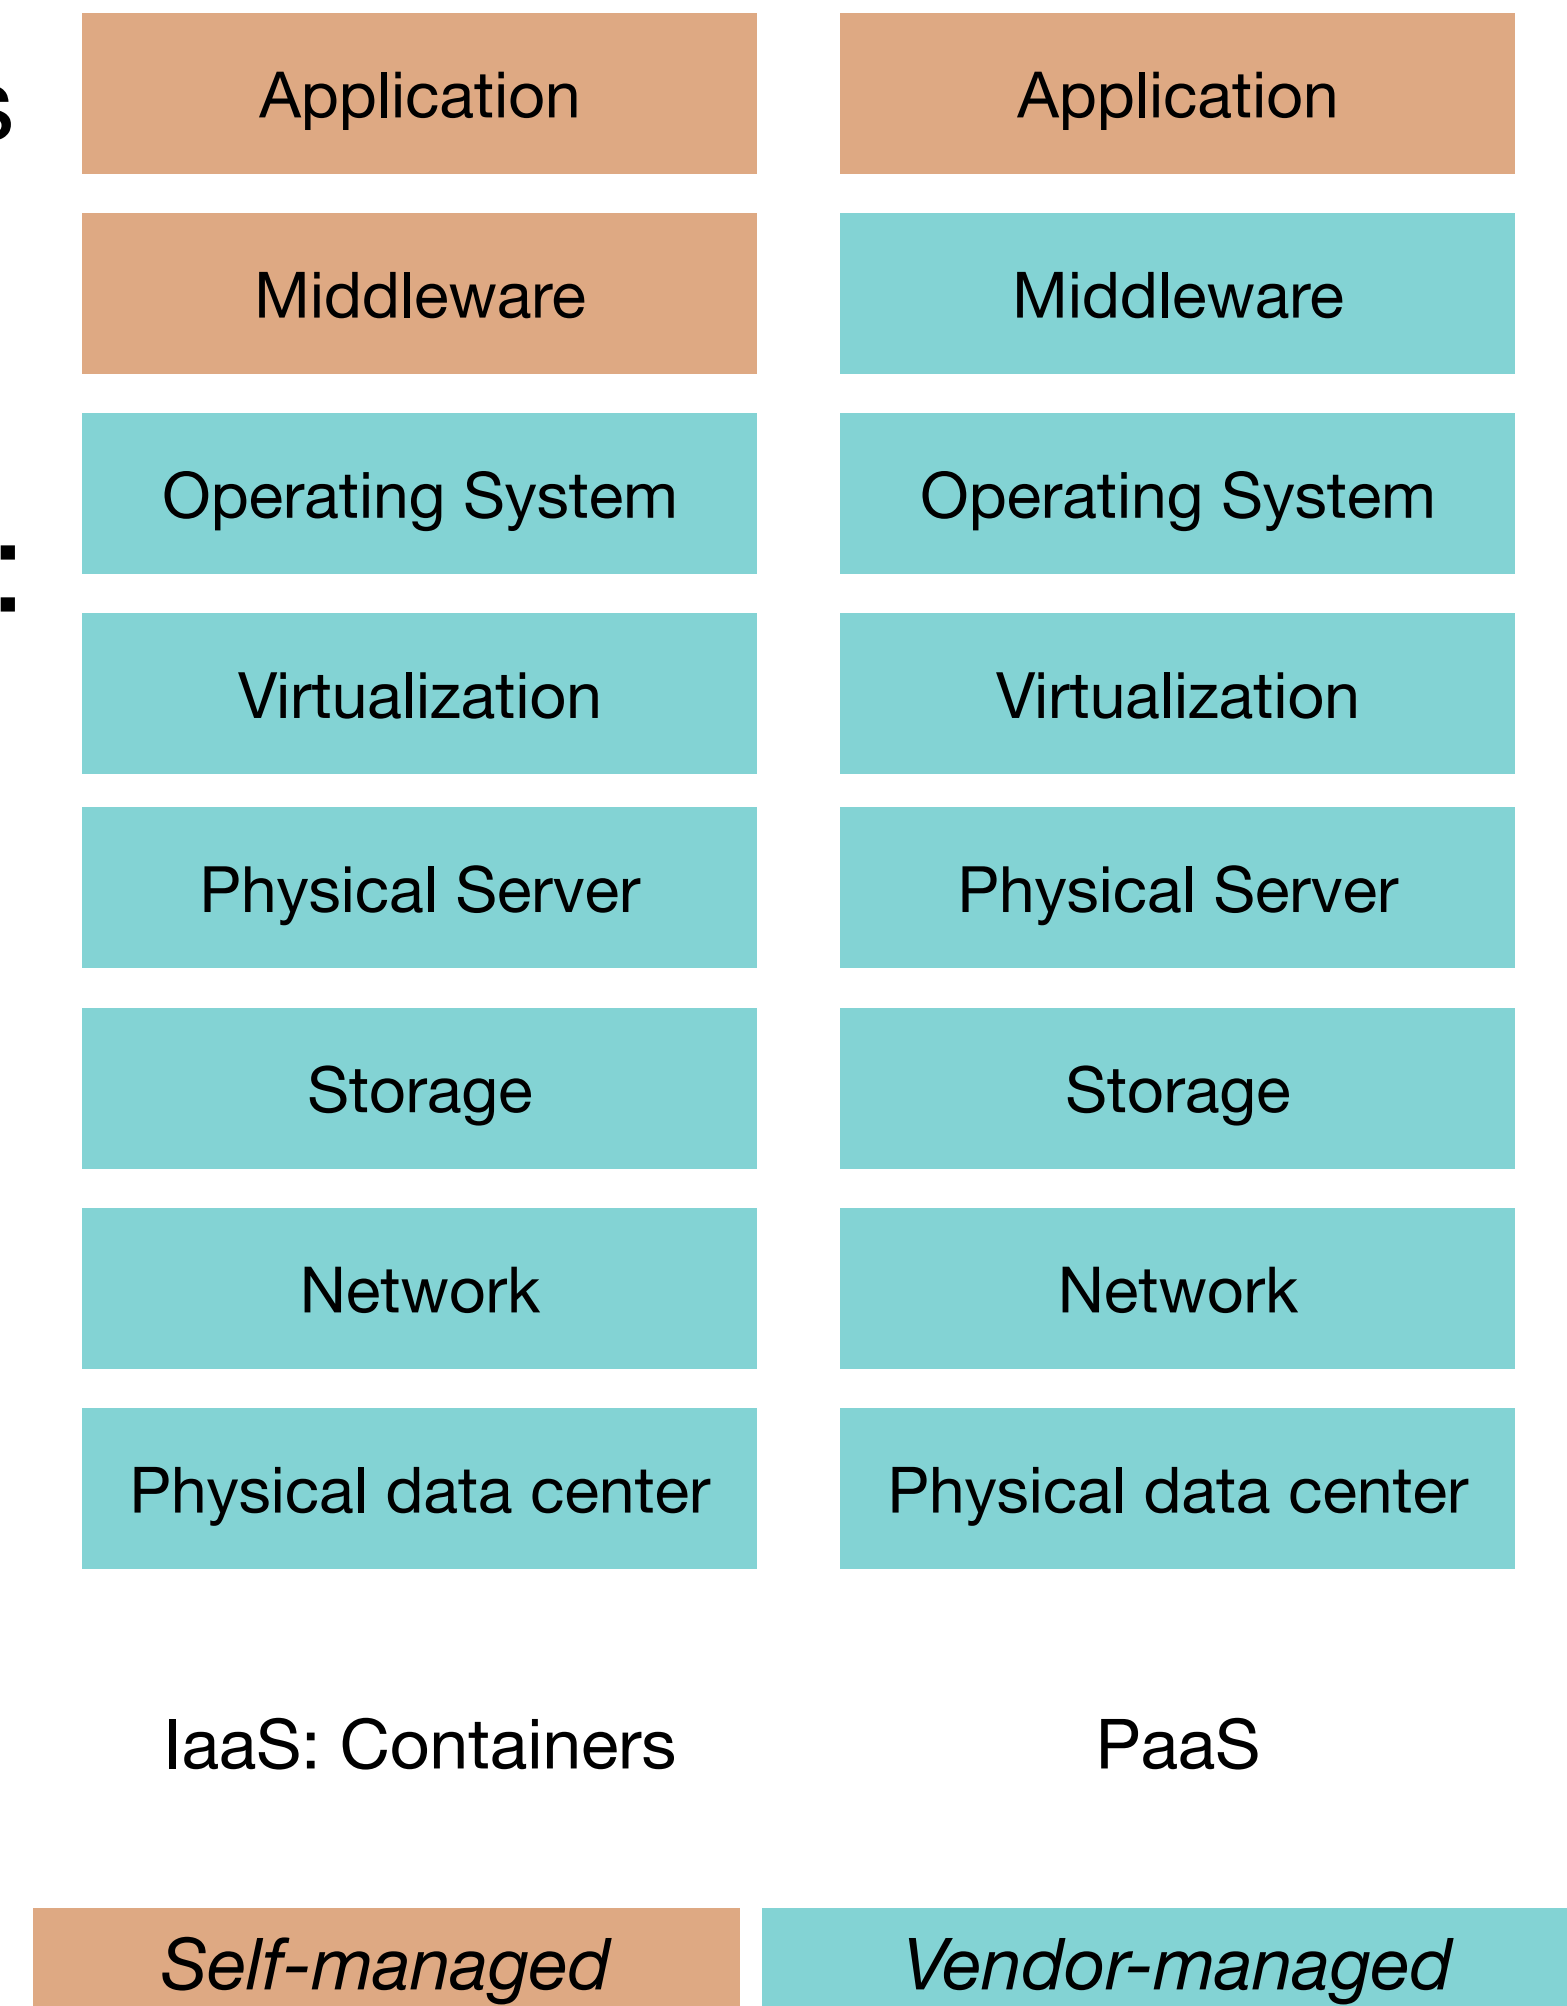


Many apps rely on common middleware

Middleware is the stuff between our app and a user's requests:

- Load balancer: route client requests to one of our app containers
- Application server: run our handler functions in response to requests from load balancer
- Monitoring/telemetry: log requests, response times and errors

Cloud vendors provide managed middleware platforms too:
“Platform as a Service”



PaaS, simplest choice for app deployment

Platform-as-a-Service provides components most apps need, fully managed by the vendor: load balancer, monitoring, application server


- Heroku, AWS Elastic Beanstalk, Google App Engine

Some PaaSs deploy apps as single functions invoked only when a web request is made

- AWS Lambda, Google Cloud Functions, Azure Functions

Some PaaSs provide databases and authentication

- Google Firebase, Back4App



PaaS

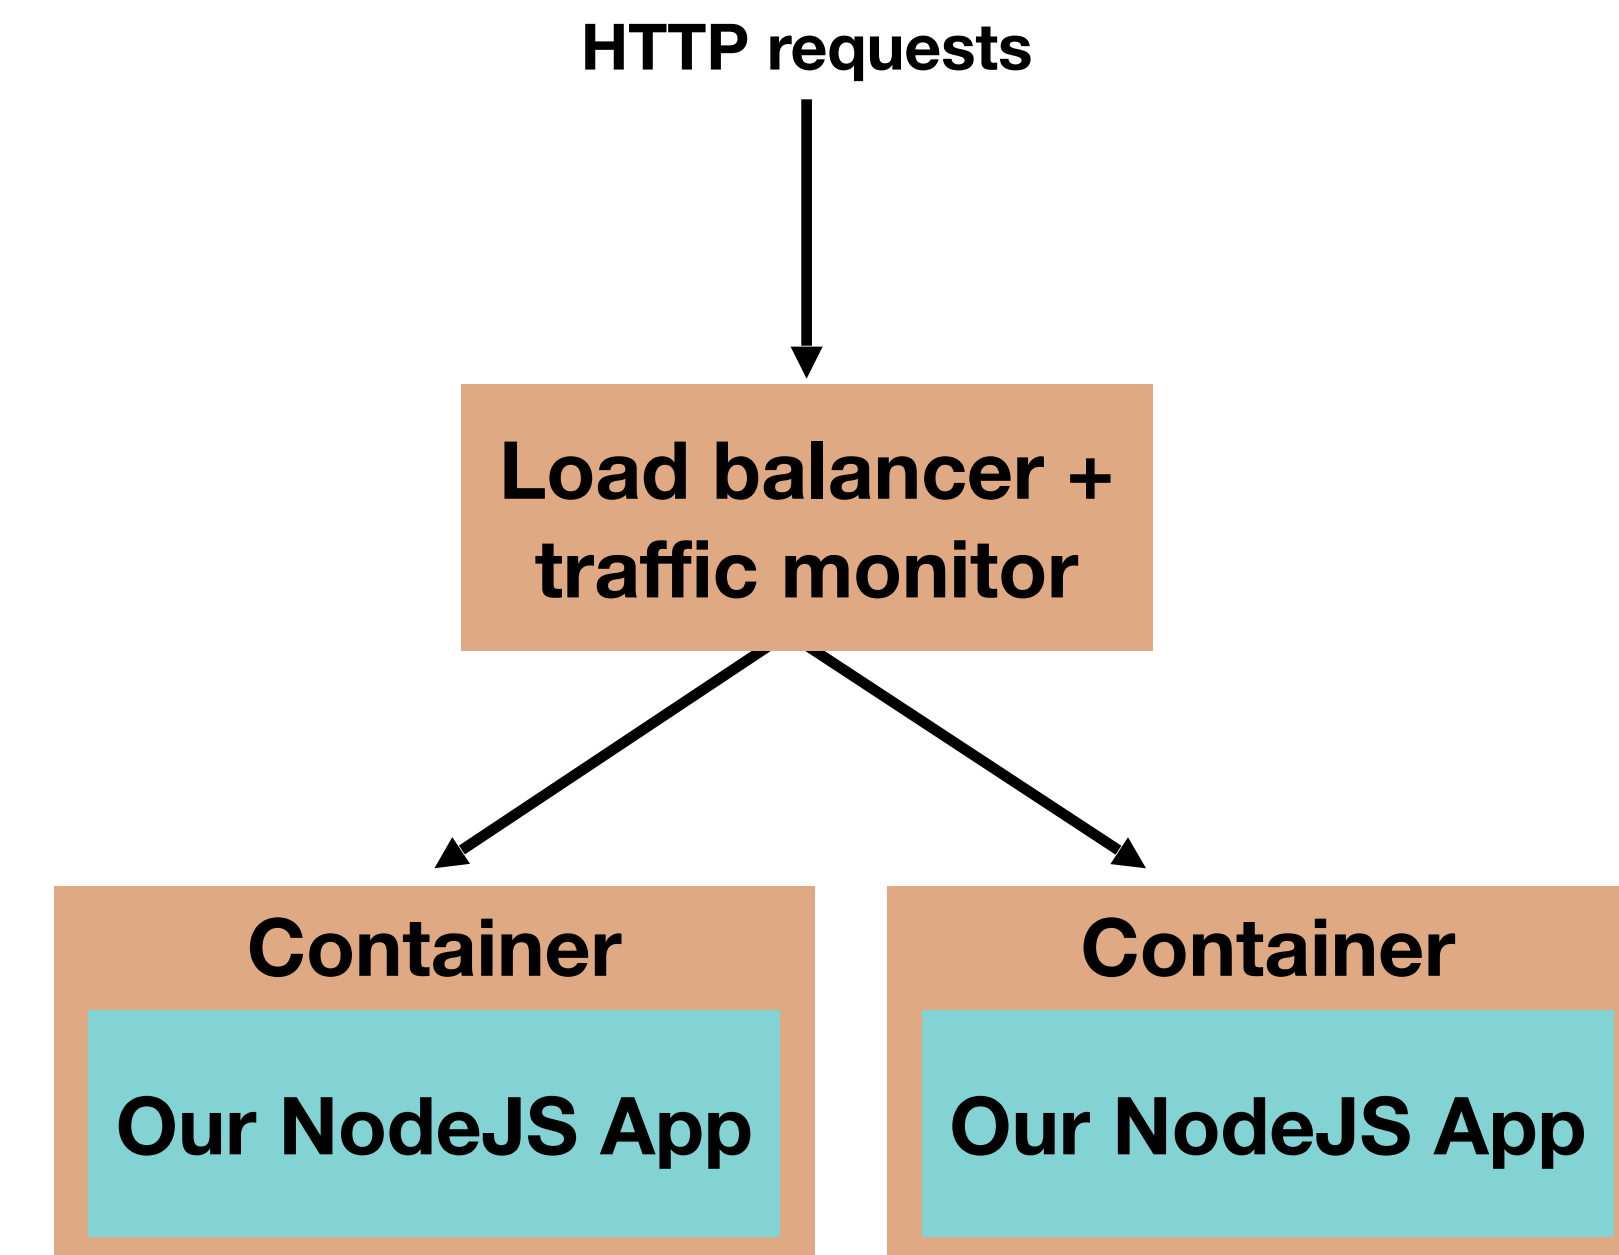
Heroku's PaaS

Takes a web app as input

- ⦿ No container, only need entry point to code, e.g. “npm start”

Hosts web app at chosen URL, can scale resources up/down on-demand

- ⦿ Load balancer fully managed by Heroku, scaling transparent
- ⦿ Auto-scale down to use no resources, spins up container on reception of a request
- ⦿ Dashboard for monitoring/reporting



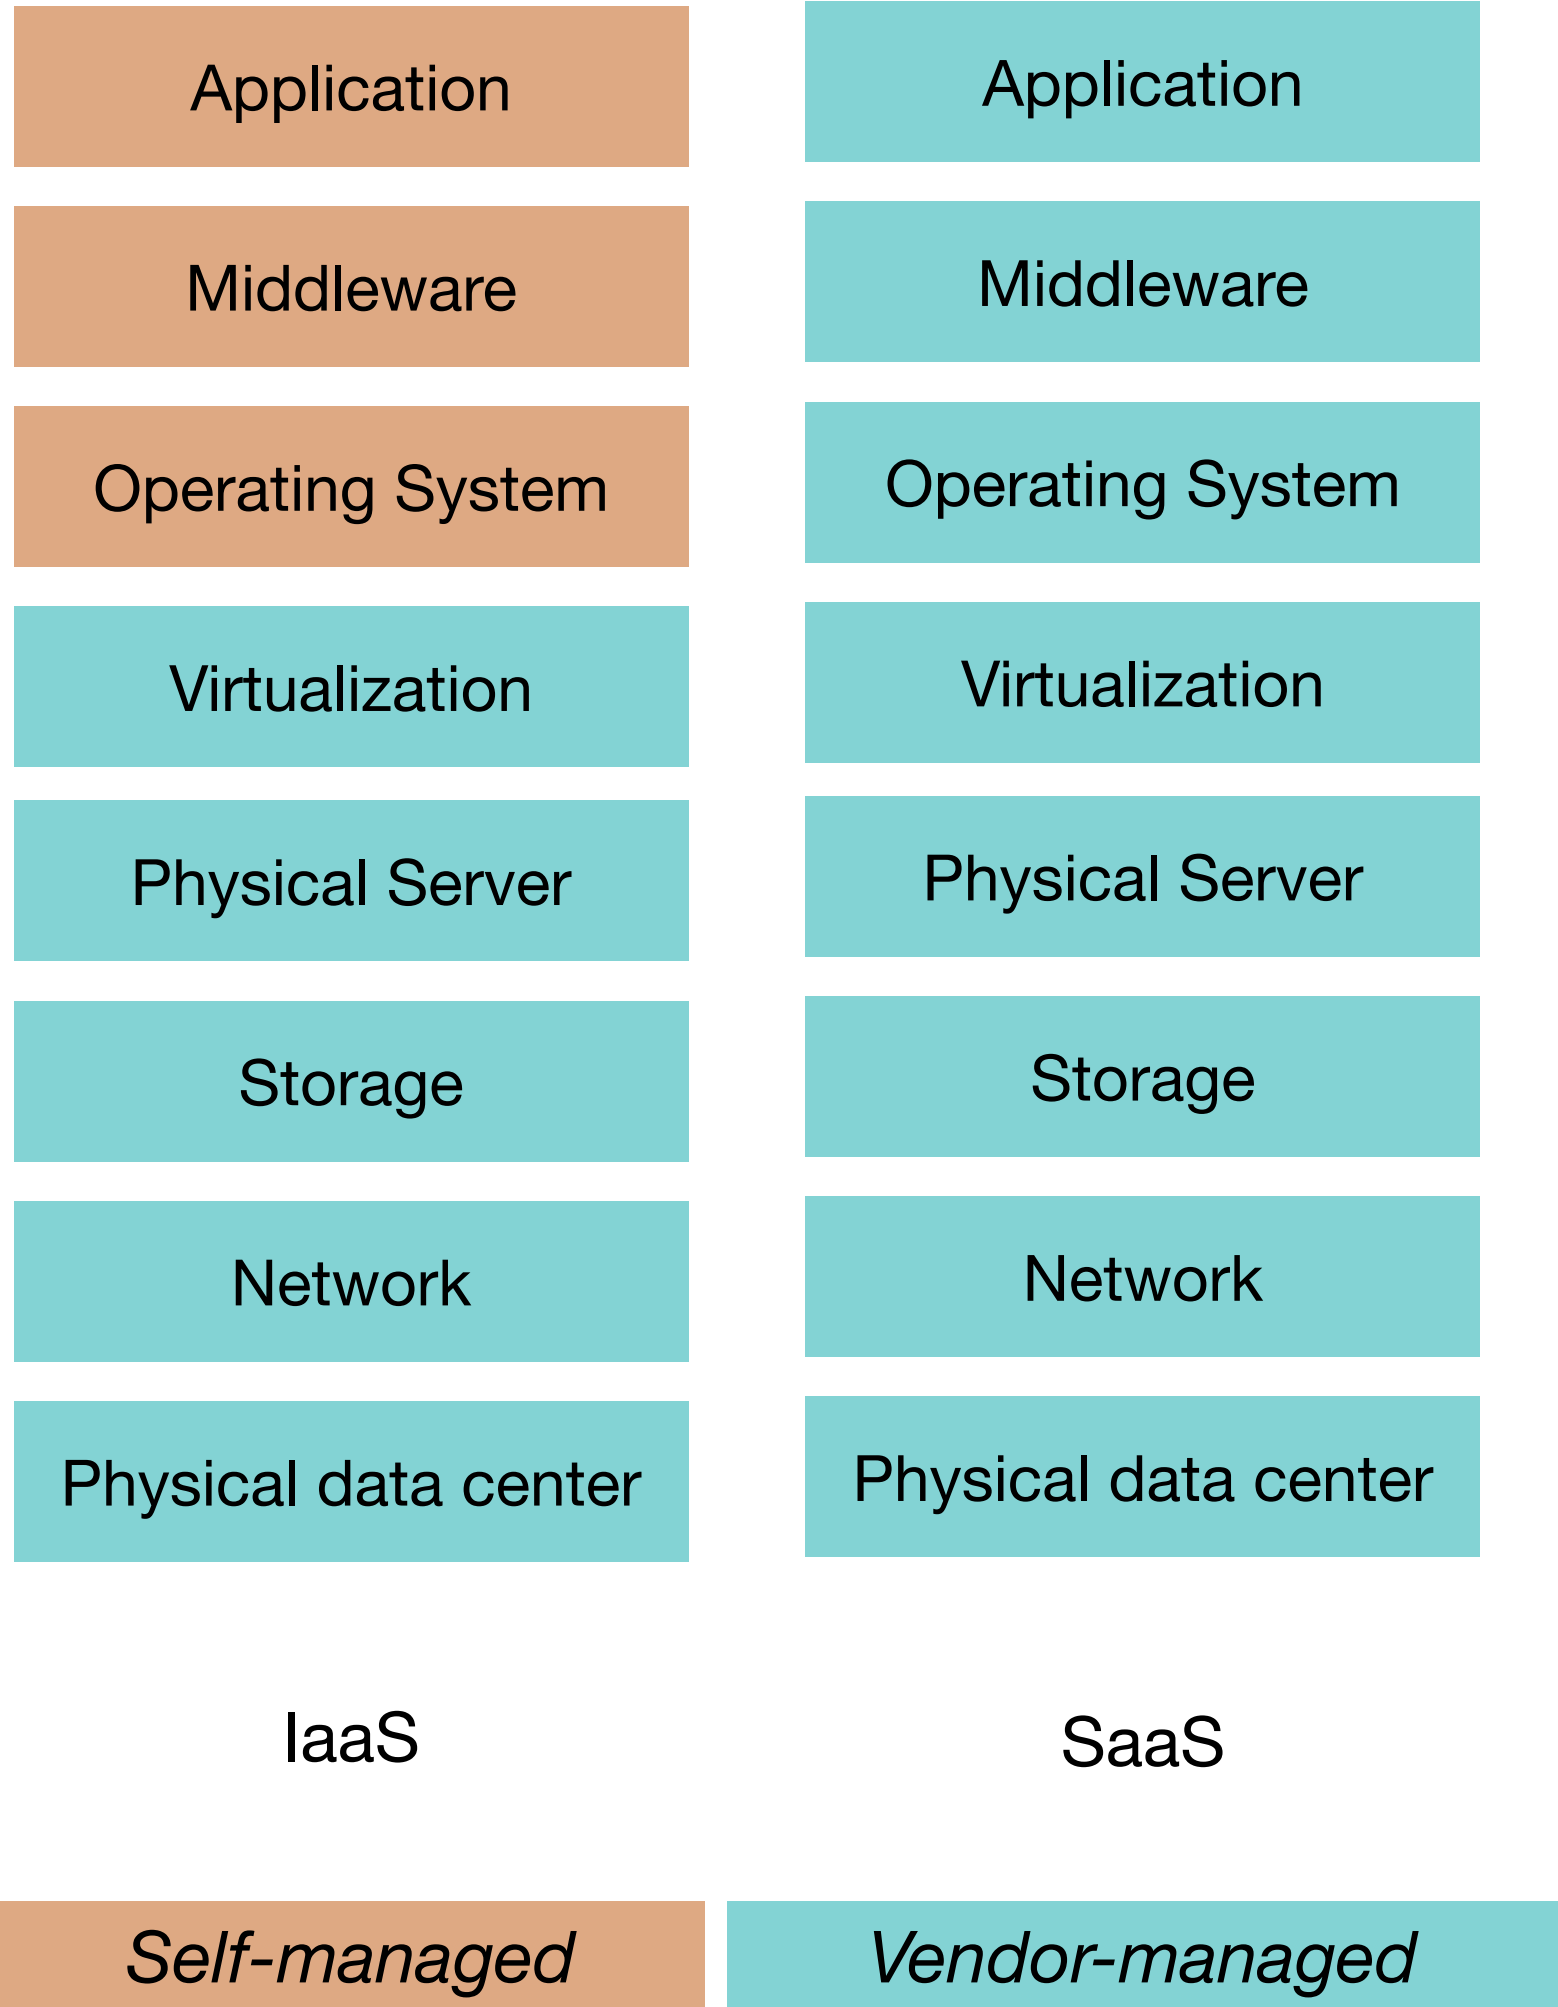
Software as a Service is fully managed

Many apps require same software, cloud providers can operate it for us

Providers also develop custom software offered only as a service

Examples:

- PostgreSQL (open source)
- Twilio Programmable Video (proprietary chat)



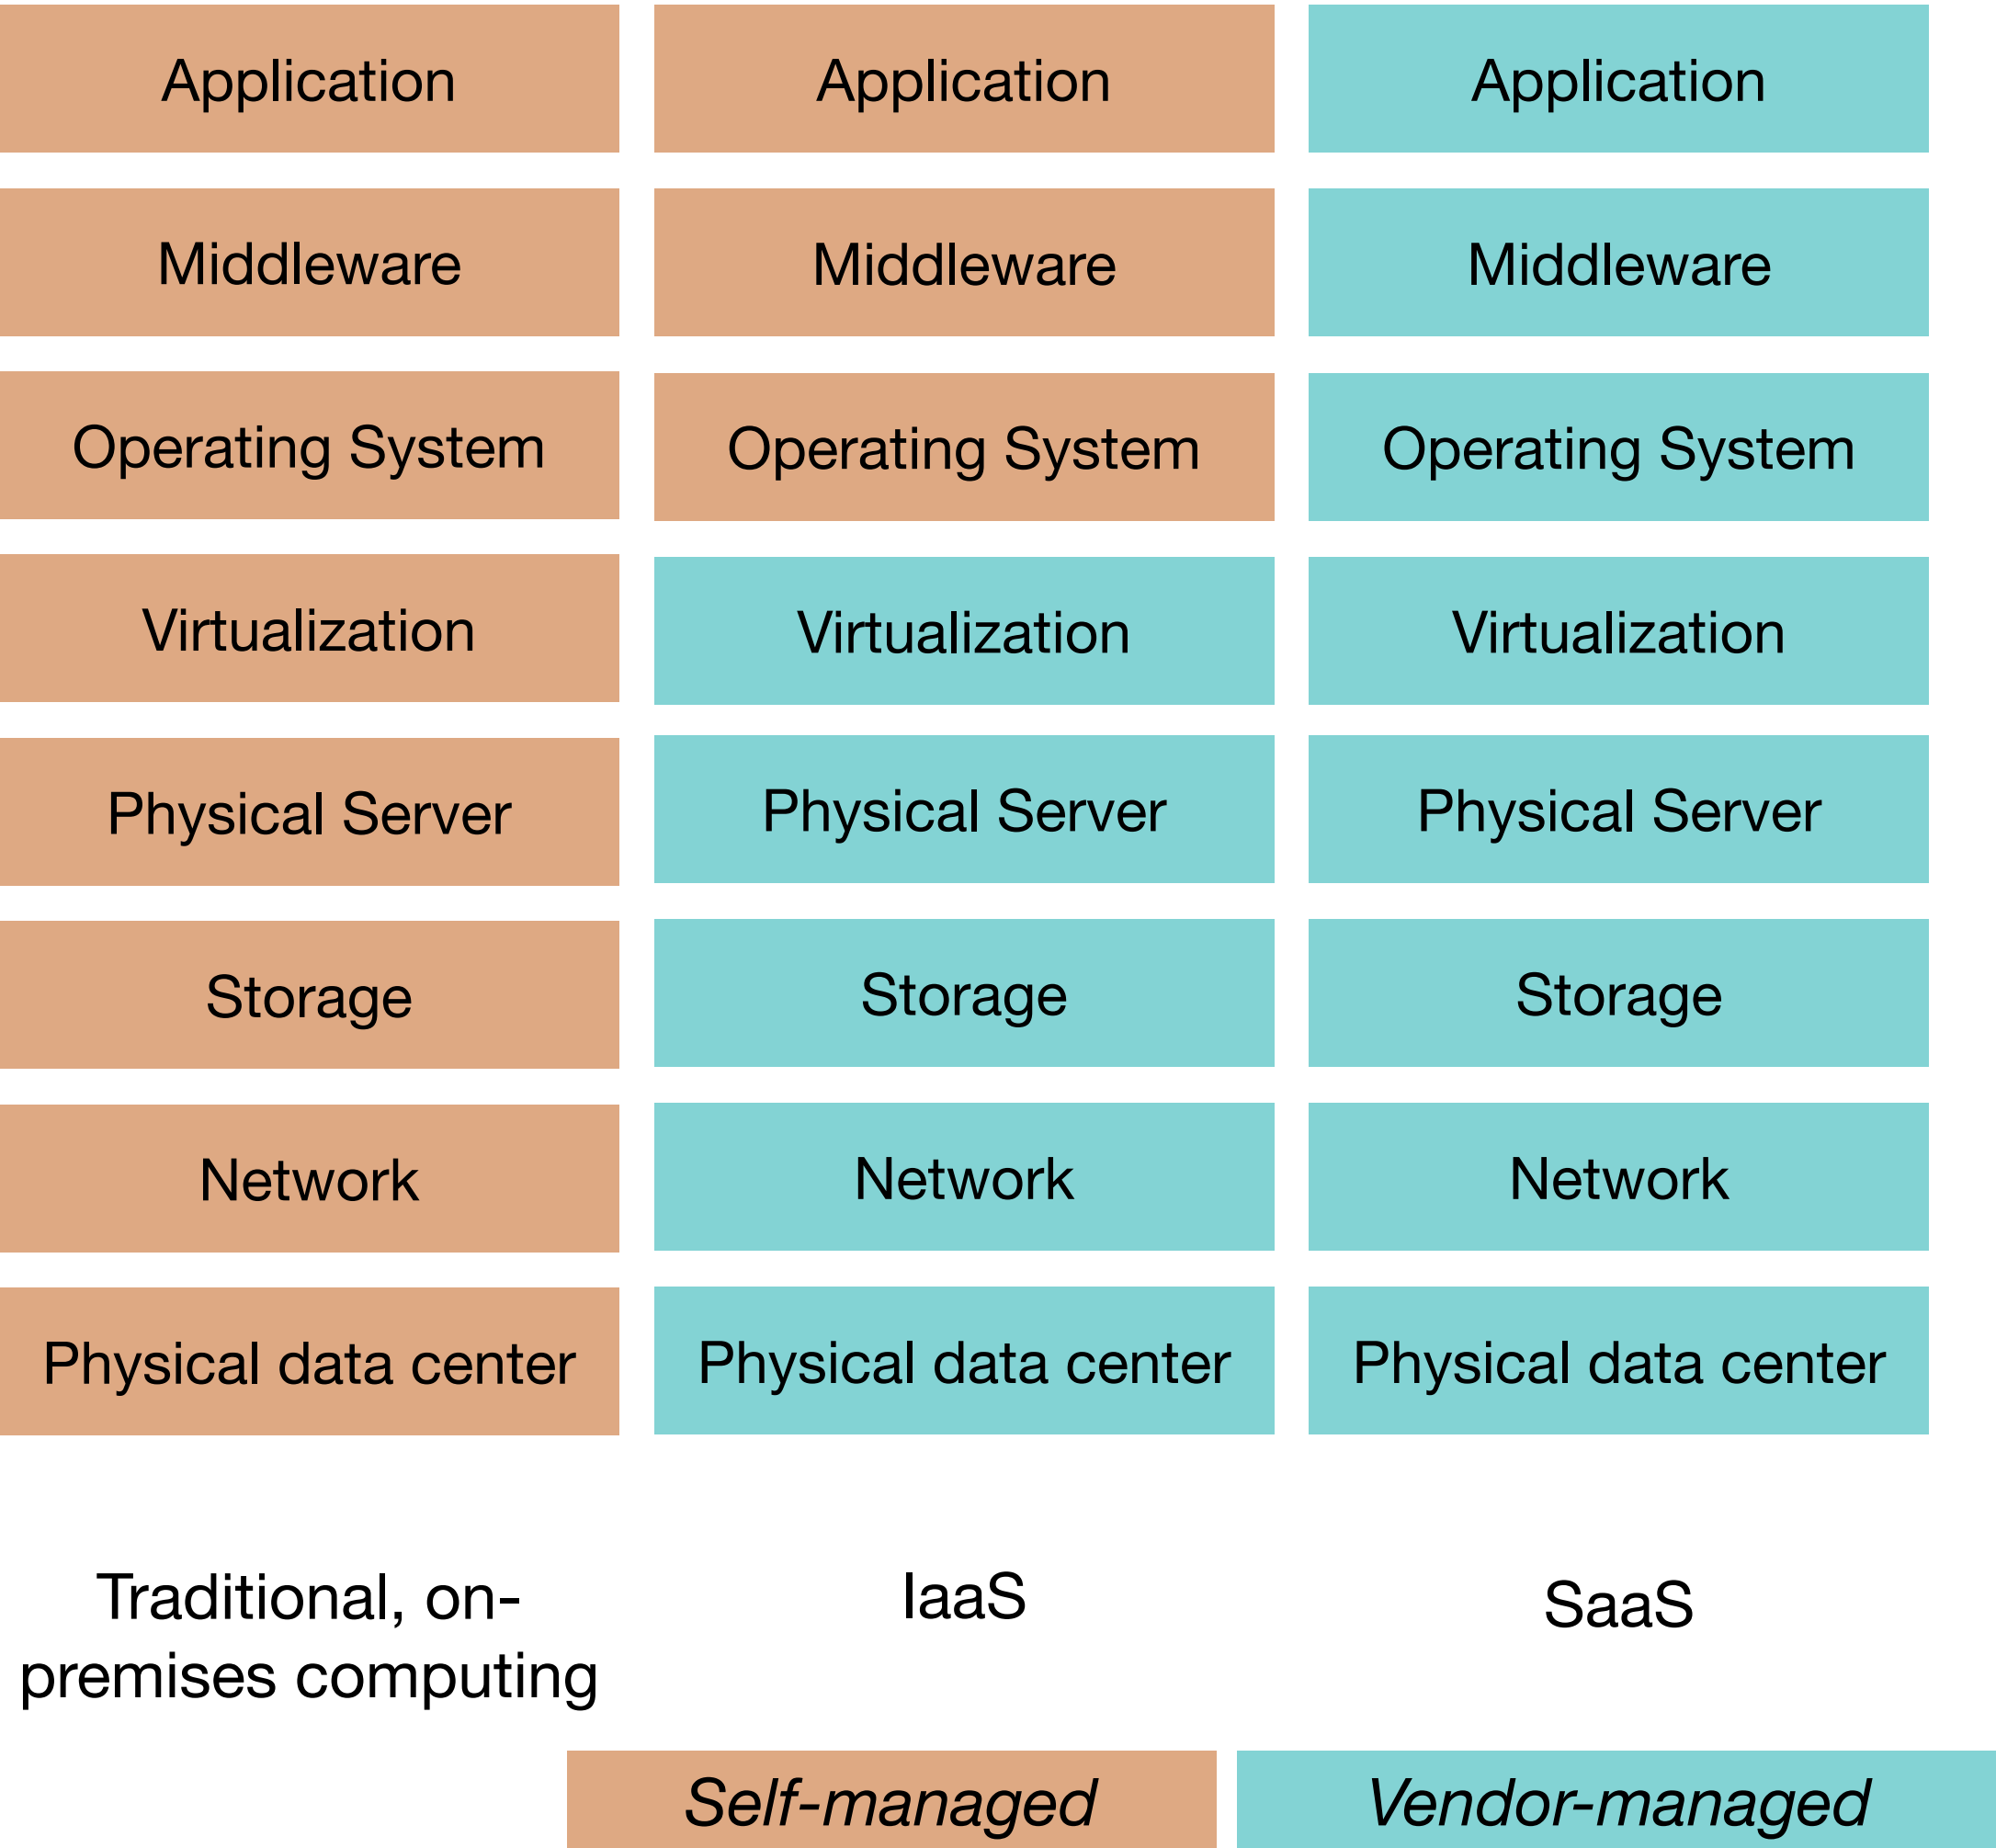
Self-managed vs Vendor-managed Infrastructure

Benefits to vendor-managed options:

- More ways to reduce resource consumption, improve resource utilization
- Less management burden
- Less capital investment, greater operating expenses

Benefits to self-managed options:

- Greater flexibility and avoid vendor lock-in
- More capital investment, less operating expenses



Cloud Infrastructure is best for variable workloads

Consider:

- ⦿ Does your workload benefit from ability to scale up or down?

Example:

- ⦿ need to run 300 VMs, each 4 vCPUs, 16GB RAM

Private cloud:

- ⦿ Dell PowerEdge Pricing (AMD EPYC 64 core CPUs)
- ⦿ 7 servers, each 128 cores, 512GB RAM, 3 TB storage = \$162,104

Public cloud:

- ⦿ Amazon EC2 Pricing (M5.xlarge instances, \$0.121/VM-hour)
- ⦿ 10 VMs for 1 year + 290 VMs for 1 month: \$36,215.30
- ⦿ 300 VMs for 1 year: \$317,988

Public clouds are not the only option

“Public” clouds are connected to the internet and available for anyone to use

- ◎ Examples: Amazon, Azure, Google Cloud, DigitalOcean

“Private” clouds use cloud technologies with on-premises, self-managed hardware

- ◎ Cost-effective when a large scale of baseline resources are needed
- ◎ Example management software: OpenStack, VMWare, Proxmox, Kubernetes

“Hybrid” clouds integrate private and public (or multiple public) clouds

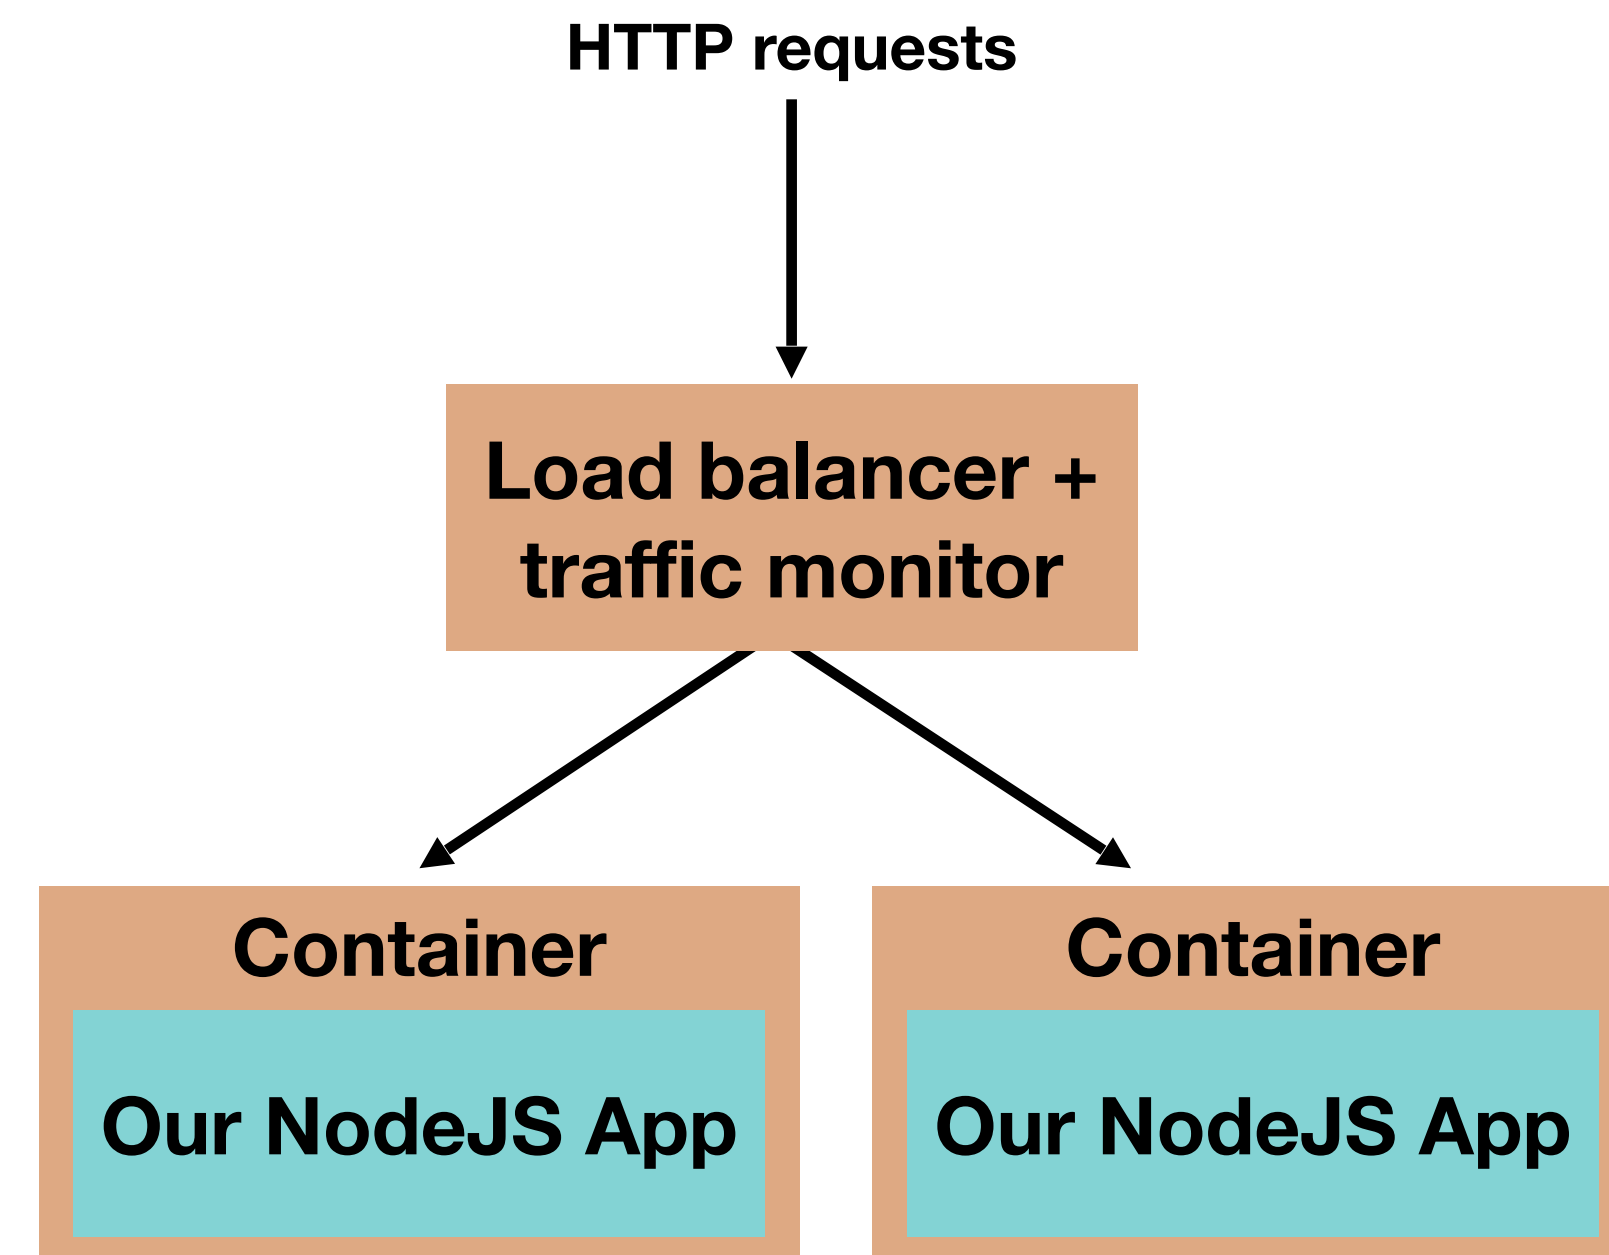
- ◎ Effective approach to “burst” capacity from private cloud to public cloud

Activity: Transcript server on Heroku

Heroku is one of the easiest platforms to use to host web apps

Objective:

- deploy our (completed) transcript server to Heroku
- Handout has detailed README



Review

You should now be able to...

- Describe what “cloud” computing is
- Understand the role of virtual machines and containers in cloud computing
- Deploy a web app to the cloud